Lecture 3: Gradient Descent Optimization

Heading: Gradient Descent
Gradient Descent is a first-order iterative optimization algorithm for finding a local minimum of a differentiable objective function. In machine learning, parameters are updated iteratively in the direction opposite to the gradient of the loss function, scaled by the learning rate. Variants include Stochastic Gradient Descent (SGD) and Mini-batch Gradient Descent.